

Stat 425 Case Study 1

Carrie Mecca, Charlie Marcou, Jack Hanley, Jessie Bustin

2022-10-19

```
# Libraries
```

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.3.2 --
## v ggplot2 3.3.6      v purrr  0.3.4
## v tibble  3.1.8      v dplyr  1.0.9
## v tidyr   1.2.0      v stringr 1.4.1
## v readr   2.1.2      v forcats 0.5.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

```
library(faraway)
```

```
library(lmtest)
```

```
## Loading required package: zoo
##
## Attaching package: 'zoo'
##
## The following objects are masked from 'package:base':
##
##      as.Date, as.Date.numeric
```

```
library(MASS)
```

```
##
## Attaching package: 'MASS'
##
## The following object is masked from 'package:dplyr':
##
##      select
```

```
# Load Data & Name Columns
```

```
data <- read_table("CDI.txt", col_names = FALSE)
```

```
##
## -- Column specification -----
## cols(
##   X1 = col_double(),
```

```
## X2 = col_character(),
## X3 = col_character(),
## X4 = col_double(),
## X5 = col_double(),
## X6 = col_double(),
## X7 = col_double(),
## X8 = col_double(),
## X9 = col_double(),
## X10 = col_double(),
## X11 = col_double(),
## X12 = col_double(),
## X13 = col_double(),
## X14 = col_double(),
## X15 = col_double(),
## X16 = col_double(),
## X17 = col_double()
## )
```

```
data <- data %>%
  rename(id = X1,
         county = X2,
         state = X3,
         land_area = X4,
         total_pop = X5,
         pop_18to24 = X6,
         pop_over65 = X7,
         num_physicians = X8,
         num_hospital_beds = X9,
         serious_crimes = X10,
         highschool_rate = X11,
         bachelors_rate = X12,
         poverty_rate = X13,
         unemployment_rate = X14,
         per_capita_income = X15,
         total_personal_income = X16,
         region = X17)
```

```
# Check Variable Types
str(data)
```

```
## spec_tbl_df [440 x 17] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ id : num [1:440] 1 2 3 4 5 6 7 8 9 10 ...
## $ county : chr [1:440] "Los_Angeles" "Cook" "Harris" "San_Diego" ...
## $ state : chr [1:440] "CA" "IL" "TX" "CA" ...
## $ land_area : num [1:440] 4060 946 1729 4205 790 ...
## $ total_pop : num [1:440] 8863164 5105067 2818199 2498016 2410556 ...
## $ pop_18to24 : num [1:440] 32.1 29.2 31.3 33.5 32.6 28.3 29.2 27.4 27.1 32.6 ...
## $ pop_over65 : num [1:440] 9.7 12.4 7.1 10.9 9.2 12.4 12.5 12.5 13.9 8.2 ...
## $ num_physicians : num [1:440] 23677 15153 7553 5905 6062 ...
## $ num_hospital_beds : num [1:440] 27700 21550 12449 6179 6369 ...
## $ serious_crimes : num [1:440] 688936 436936 253526 173821 144524 ...
## $ highschool_rate : num [1:440] 70 73.4 74.9 81.9 81.2 63.7 81.5 70 65 77.1 ...
## $ bachelors_rate : num [1:440] 22.3 22.8 25.4 25.3 27.8 16.6 22.1 13.7 18.8 26.3 ...
```

```
## $ poverty_rate      : num [1:440] 11.6 11.1 12.5 8.1 5.2 19.5 8.8 16.9 14.2 10.4 ...
## $ unemployment_rate : num [1:440] 8 7.2 5.7 6.1 4.8 9.5 4.9 10 8.7 6.1 ...
## $ per_capita_income  : num [1:440] 20786 21729 19517 19588 24400 ...
## $ total_personal_income: num [1:440] 184230 110928 55003 48931 58818 ...
## $ region             : num [1:440] 4 2 3 4 4 1 4 2 3 3 ...
## - attr(*, "spec")=
## .. cols(
## ..   X1 = col_double(),
## ..   X2 = col_character(),
## ..   X3 = col_character(),
## ..   X4 = col_double(),
## ..   X5 = col_double(),
## ..   X6 = col_double(),
## ..   X7 = col_double(),
## ..   X8 = col_double(),
## ..   X9 = col_double(),
## ..   X10 = col_double(),
## ..   X11 = col_double(),
## ..   X12 = col_double(),
## ..   X13 = col_double(),
## ..   X14 = col_double(),
## ..   X15 = col_double(),
## ..   X16 = col_double(),
## ..   X17 = col_double()
## .. )
```

```
# Check for NA's and INF's
complete_rows <- data[complete.cases(data), ]
nrow(data) == nrow(complete_rows)
```

```
## [1] TRUE
```

```
# Change region type because it is not numeric
data <- data %>%
  mutate(region = as.factor(region))

# Check correlation for numeric features
data %>%
  dplyr::select(-id, -county, -state, -region) %>%
  cor() %>%
  round(digits = 2)
```

```
##               land_area total_pop pop_18to24 pop_over65 num_physicians
## land_area          1.00      0.17      -0.05       0.01         0.08
## total_pop           0.17      1.00       0.08      -0.03         0.94
## pop_18to24         -0.05      0.08       1.00      -0.62         0.12
## pop_over65          0.01     -0.03      -0.62       1.00         0.00
## num_physicians      0.08      0.94       0.12       0.00         1.00
## num_hospital_beds   0.07      0.92       0.07       0.05         0.95
## serious_crimes      0.13      0.89       0.09      -0.04         0.82
## highschool_rate     -0.10     -0.02       0.25      -0.27         0.00
## bachelors_rate      -0.14      0.15       0.46      -0.34         0.24
## poverty_rate        0.17      0.04       0.03       0.01         0.06
```

```

## unemployment_rate      0.20      0.01     -0.28      0.24      -0.05
## per_capita_income      -0.19      0.24     -0.03      0.02       0.32
## total_personal_income   0.13      0.99      0.07     -0.02       0.95
##                          num_hospital_beds serious_crimes highschool_rate
## land_area              0.07          0.13          -0.10
## total_pop              0.92          0.89          -0.02
## pop_18to24             0.07          0.09           0.25
## pop_over65             0.05         -0.04          -0.27
## num_physicians         0.95          0.82           0.00
## num_hospital_beds      1.00          0.86          -0.11
## serious_crimes         0.86          1.00          -0.11
## highschool_rate       -0.11         -0.11           1.00
## bachelors_rate         0.10          0.08           0.71
## poverty_rate           0.17          0.16          -0.69
## unemployment_rate      0.01          0.04          -0.59
## per_capita_income       0.19          0.12           0.52
## total_personal_income   0.90          0.84           0.04
##                          bachelors_rate poverty_rate unemployment_rate
## land_area             -0.14          0.17           0.20
## total_pop              0.15          0.04           0.01
## pop_18to24             0.46          0.03          -0.28
## pop_over65            -0.34          0.01           0.24
## num_physicians         0.24          0.06          -0.05
## num_hospital_beds      0.10          0.17           0.01
## serious_crimes         0.08          0.16           0.04
## highschool_rate        0.71         -0.69          -0.59
## bachelors_rate         1.00         -0.41          -0.54
## poverty_rate          -0.41          1.00           0.44
## unemployment_rate     -0.54          0.44           1.00
## per_capita_income       0.70         -0.60          -0.32
## total_personal_income   0.22         -0.04          -0.03
##                          per_capita_income total_personal_income
## land_area             -0.19          0.13
## total_pop              0.24          0.99
## pop_18to24            -0.03          0.07
## pop_over65             0.02         -0.02
## num_physicians         0.32          0.95
## num_hospital_beds      0.19          0.90
## serious_crimes         0.12          0.84
## highschool_rate        0.52          0.04
## bachelors_rate         0.70          0.22
## poverty_rate          -0.60         -0.04
## unemployment_rate     -0.32         -0.03
## per_capita_income       1.00          0.35
## total_personal_income   0.35          1.00

```

```

# We have 4 variables highly correlated with total_pop so will transform them to per 100,000 people for
# Will drop total income because we already have per capita
# We can also drop the ID column

```

```

model_data <- data.frame(data %>%
  dplyr::select(-id, -total_personal_income) %>%
  mutate(hospital_beds_percap = num_hospital_beds / (total_pop / 100000),
         serious_crimes_percap = serious_crimes / (total_pop / 100000)) %>%
  dplyr::select(-num_hospital_beds, -serious_crimes))

```

```
# Check county levels
model_data %>%
  group_by(county) %>%
  summarise(counts = n()) %>%
  summarise(min = min(counts), max(counts))
```

```
## # A tibble: 1 x 2
##   min 'max(counts)'
##   <int>         <int>
## 1     1           7
```

```
# Drop Counties and check states
model_data <- model_data %>%
  dplyr::select(-county)
model_data %>%
  group_by(state) %>%
  summarise(counts = n()) %>%
  arrange(counts) %>%
  head()
```

```
## # A tibble: 6 x 2
##   state counts
##   <chr>   <int>
## 1 DC         1
## 2 ID         1
## 3 MT         1
## 4 ND         1
## 5 SD         1
## 6 VT         1
```

```
# Drop States and we will use regions
model_data <- model_data %>%
  dplyr::select(-state)

# Recheck correlation
model_data %>%
  dplyr::select(-region) %>%
  cor() %>%
  round(digits = 2)
```

```
##           land_area total_pop pop_18to24 pop_over65 num_physicians
## land_area         1.00    0.17   -0.05    0.01         0.08
## total_pop         0.17    1.00    0.08   -0.03         0.94
## pop_18to24       -0.05    0.08    1.00   -0.62         0.12
## pop_over65        0.01   -0.03   -0.62    1.00         0.00
## num_physicians    0.08    0.94    0.12    0.00         1.00
## highschool_rate  -0.10   -0.02    0.25   -0.27         0.00
## bachelors_rate   -0.14    0.15    0.46   -0.34         0.24
## poverty_rate      0.17    0.04    0.03    0.01         0.06
## unemployment_rate 0.20    0.01   -0.28    0.24        -0.05
## per_capita_income -0.19    0.24   -0.03    0.02         0.32
```

```

## hospital_beds_percap      -0.14      0.02      0.03      0.25      0.19
## serious_crimes_percap     0.04      0.28      0.19     -0.07      0.31
##                          highschool_rate bachelors_rate poverty_rate
## land_area                -0.10                -0.14      0.17
## total_pop                -0.02                0.15      0.04
## pop_18to24               0.25                0.46      0.03
## pop_over65              -0.27               -0.34      0.01
## num_physicians           0.00                0.24      0.06
## highschool_rate          1.00                0.71     -0.69
## bachelors_rate           0.71                1.00     -0.41
## poverty_rate            -0.69               -0.41      1.00
## unemployment_rate        -0.59               -0.54      0.44
## per_capita_income         0.52                0.70     -0.60
## hospital_beds_percap     -0.21               -0.05      0.37
## serious_crimes_percap    -0.23                0.04      0.47
##                          unemployment_rate per_capita_income hospital_beds_percap
## land_area                0.20                -0.19     -0.14
## total_pop                0.01                0.24      0.02
## pop_18to24              -0.28               -0.03      0.03
## pop_over65              0.24                0.02      0.25
## num_physicians          -0.05                0.32      0.19
## highschool_rate         -0.59                0.52     -0.21
## bachelors_rate          -0.54                0.70     -0.05
## poverty_rate            0.44               -0.60      0.37
## unemployment_rate        1.00               -0.32     -0.06
## per_capita_income        -0.32                1.00     -0.05
## hospital_beds_percap     -0.06               -0.05      1.00
## serious_crimes_percap     0.04               -0.08      0.36
##                          serious_crimes_percap
## land_area                0.04
## total_pop                0.28
## pop_18to24               0.19
## pop_over65              -0.07
## num_physicians           0.31
## highschool_rate         -0.23
## bachelors_rate           0.04
## poverty_rate            0.47
## unemployment_rate        0.04
## per_capita_income        -0.08
## hospital_beds_percap     0.36
## serious_crimes_percap     1.00

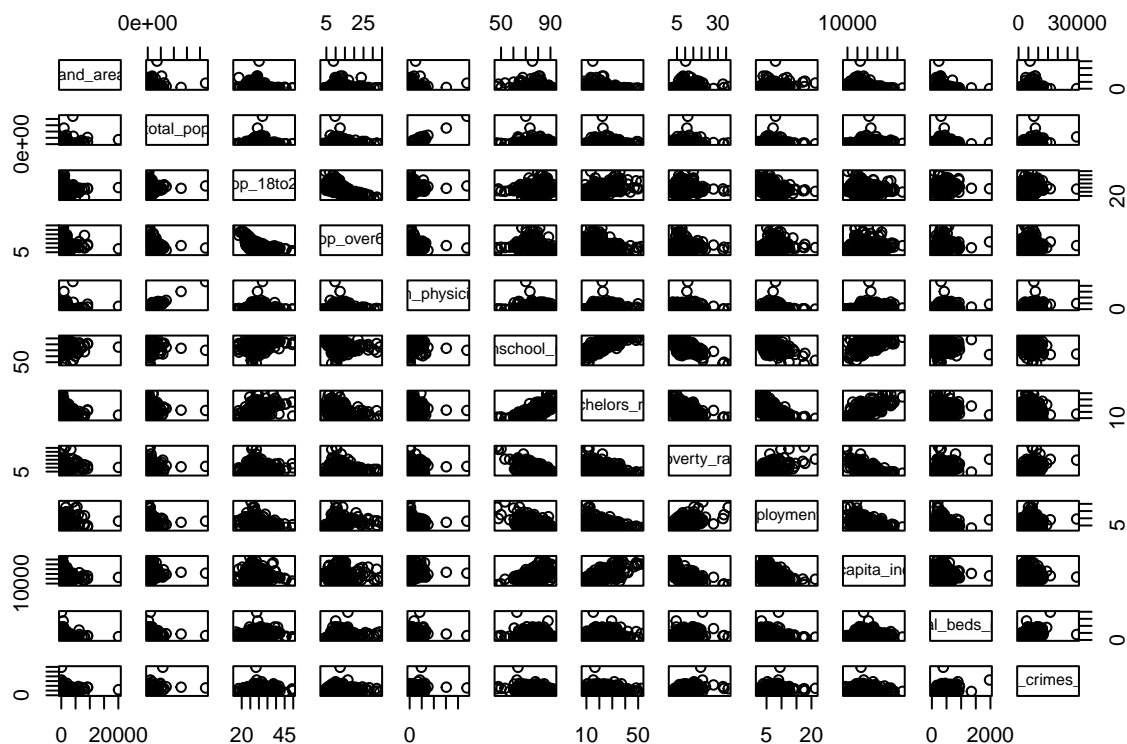
```

```
## Let's also look at a scattermatrix for correlation
```

```

model_data %>%
  dplyr::select(-region) %>%
  pairs()

```



With a cutoff of 0.75 we have no highly correlated pairs other than population and our target

With a cutoff of 0.7 we have the 2 following pairs:

high_school_rate & bachelors_rate

bachelors_rate & per_capita_income

Start by just looking at total_pop for fun

```
slr_model <- lm(num_physicians~total_pop, data = model_data)
summary(slr_model)
```

```
##
## Call:
## lm(formula = num_physicians ~ total_pop, data = model_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1969.4  -209.2   -88.0    27.9   3928.7
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.106e+02  3.475e+01  -3.184  0.00156 **
## total_pop    2.795e-03  4.837e-05  57.793  < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

```
## Residual standard error: 610.1 on 438 degrees of freedom
## Multiple R-squared:  0.8841, Adjusted R-squared:  0.8838
## F-statistic: 3340 on 1 and 438 DF,  p-value: < 2.2e-16
```

Full Model

```
mlr_full_model <- lm(num_physicians~., data = model_data)
summary(mlr_full_model)
```

```
##
## Call:
## lm(formula = num_physicians ~ ., data = model_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1904.06  -241.05   -30.21   185.69  2702.83
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -1.284e+03  7.154e+02  -1.795  0.073363 .
## land_area      -6.667e-02  1.734e-02  -3.844  0.000139 ***
## total_pop       2.708e-03  4.194e-05  64.571 < 2e-16 ***
## pop_18to24      1.860e+01  8.558e+00   2.174  0.030266 *
## pop_over65      8.983e+00  7.922e+00   1.134  0.257460
## highschool_rate -1.372e+01  6.937e+00  -1.978  0.048613 *
## bachelors_rate  1.549e+01  7.499e+00   2.065  0.039493 *
## poverty_rate    2.493e+01  1.030e+01   2.421  0.015875 *
## unemployment_rate -1.479e+01  1.381e+01  -1.071  0.284764
## per_capita_income  4.812e-02  1.218e-02   3.952  9.07e-05 ***
## region2        -3.810e+01  7.129e+01  -0.534  0.593340
## region3        -6.191e+01  7.379e+01  -0.839  0.401954
## region4         1.700e+02  8.910e+01   1.908  0.057117 .
## hospital_beds_percap  1.448e+00  1.511e-01   9.582 < 2e-16 ***
## serious_crimes_percap -2.903e-02  1.117e-02  -2.599  0.009668 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 462.3 on 425 degrees of freedom
## Multiple R-squared:  0.9354, Adjusted R-squared:  0.9333
## F-statistic: 439.6 on 14 and 425 DF,  p-value: < 2.2e-16
```

Remove some fields based on EDA and full model summary

```
sub1_data <- model_data %>%
  dplyr::select(-pop_over65, -unemployment_rate, -highschool_rate)
mlr_sub1_model <- lm(num_physicians~., data = sub1_data)
anova(mlr_sub1_model, mlr_full_model)
```

Analysis of Variance Table

```
##
## Model 1: num_physicians ~ land_area + total_pop + pop_18to24 + bachelors_rate +
##      poverty_rate + per_capita_income + region + hospital_beds_percap +
##      serious_crimes_percap
## Model 2: num_physicians ~ land_area + total_pop + pop_18to24 + pop_over65 +
##      highschool_rate + bachelors_rate + poverty_rate + unemployment_rate +
```



```
##      per_capita_income + region + hospital_beds_percap + serious_crimes_percap
## Res.Df      RSS Df Sum of Sq      F Pr(>F)
## 1      428 92073649
## 2      425 90825598  3   1248050 1.9467 0.1214
```


```
## Checking high-leverage points
```

```
leverages=lm.influence(mlr_sub1_model)$hat
head(leverages)
```

```
##           1           2           3           4           5           6
## 0.51964431 0.16519116 0.04881486 0.04305255 0.05235414 0.31889087
```

```
## Plot to help identify high leverage observations
```

```
halfnorm(leverages, nlab=6, labs=as.character(1:length(leverages)), ylab="Leverages")
```



```
## Determining leverages that exceed a 2p/n threshold
```

```
n = dim(model_data)[1]
p = length(variable.names(mlr_sub1_model))
leverages.high = leverages[leverages>(2*p/n)]
leverages.high
```

```
##           1           2           6           7           14           42           48
## 0.51964431 0.16519116 0.31889087 0.08057521 0.41713279 0.05953933 0.05626570
##           49           65           67           85           95          123          128
```

```
## 0.06403587 0.07669668 0.06125503 0.05643133 0.05796254 0.16774044 0.14336932
##      187      188      206      235      303      337      357
## 0.05963580 0.11180418 0.09706308 0.06618486 0.11559560 0.10010948 0.05624120
##      363      392      396      400      405      412      418
## 0.07263222 0.05937063 0.06016817 0.06739193 0.06502508 0.08302403 0.09746487
##      433
## 0.05524000
```

```
## We currently have many high leverage points (29), They represent only about 6.6% of the data.
## Before continuing, let us look at what high leverage points are good and bad
## Calculate IQR for number of physicians
IQR_y = IQR(model_data$num_physicians)
## Define range with its lower limit being (Q1 - IQR) and upper limit being (Q3 + IQR)
QT1_y = quantile(model_data$num_physicians,0.25)
QT3_y = quantile(model_data$num_physicians,0.75)
lower_lim_y = QT1_y - IQR_y
upper_lim_y = QT3_y + IQR_y
vector_lim_y = c(lower_lim_y,upper_lim_y)
## Range for number of physicians
vector_lim_y
```

```
##      25%      75%
## -670.50 1889.25
```

```
## Extract observations with high leverage points from the original data frame
highlev = model_data[leverages>2*p/n,]
## Select only the observations with leverage points outside the range
highlev_lower = highlev[highlev$num_physicians < vector_lim_y[1], ]
highlev_upper = highlev[highlev$num_physicians > vector_lim_y[2], ]
highlev2 = rbind(highlev_lower,highlev_upper)
## This is not outputting the observation number like her example did. It is probably because we're using
##I switched model_data to be a dataframe which I believe solves this issue-Carrie
highlev2
```

```
##      land_area total_pop pop_18to24 pop_over65 num_physicians highschool_rate
## 1      4060      8863164      32.1      9.7      23677      70.0
## 2      946      5105067      29.2      12.4      15153      73.4
## 6      71      2300664      28.3      12.4      4861      63.7
## 7      9204      2122101      29.2      12.5      4320      81.5
## 14      20062      1418380      30.1      8.8      2463      75.4
## 48      495      757027      28.6      10.2      4635      90.6
## 67      59      663906      39.2      12.1      5674      75.4
## 95      181      496938      28.3      13.0      2500      68.1
## 123      62      396685      28.7      16.6      4189      62.8
##      bachelors_rate poverty_rate unemployment_rate per_capita_income region
## 1      22.3      11.6      8.0      20786      4
## 2      22.8      11.1      7.2      21729      2
## 6      16.6      19.5      9.5      16803      1
## 7      22.1      8.8      4.9      18042      4
## 14      14.9      10.3      8.0      16399      4
## 48      49.9      2.7      3.3      30081      3
## 67      27.7      14.4      8.7      23150      1
## 95      22.4      27.3      6.1      16578      3
```

```
## 123          15.3          20.6          9.0          18113          2
##      hospital_beds_percap serious_crimes_percap
## 1          312.5295          7773.026
## 2          422.1296          8558.869
## 6          388.6704          29598.672
## 7          287.6395          8368.735
## 14         236.1144          5859.502
## 48         199.0682          4590.853
## 67         926.9385          10364.118
## 95         808.5516          10914.440
## 123        1969.8249          16159.673
```

```
## Computing Studentized Residuals
mlr_sub1_model$resid = rstudent(mlr_sub1_model);
## Critical value with Bonferroni correction
## Note: Compare to t-value later at the alpha we choose
bonferroni_cv = qt(.05/(2*n), n-p-1)
bonferroni_cv
```

```
## [1] -3.895681
```

```
## Sorting residuals to find outliers
mlr_sub1_model$resid.sorted = sort(abs(mlr_sub1_model$resid), decreasing=TRUE)[1:10]
print(mlr_sub1_model$resid.sorted)
```

```
##      50      67      48      19      8      53      11      32
## 6.267323 5.790836 5.067090 4.513103 4.376934 4.185103 4.059736 3.088757
##      15      2
## 3.059593 2.722946
```


```
## Printing those above the value
## We can see observations 50, 67, 48, 19, 8, 53, and 11 are outliers.
mlr_sub1_model$outliers = mlr_sub1_model$resid.sorted[abs(mlr_sub1_model$resid.sorted) > abs(bonferroni_cv)]
print(mlr_sub1_model$outliers)
```

```
##      50      67      48      19      8      53      11
## 6.267323 5.790836 5.067090 4.513103 4.376934 4.185103 4.059736
```

```
## Finding high cook's distance observations
mlr_sub1_model$cooks = cooks.distance(mlr_sub1_model)
sort(mlr_sub1_model$cooks, decreasing = TRUE)[1:10]
```


```
##      6      67      1      48      2      50      53
## 0.22208251 0.16946408 0.15123922 0.12061098 0.12045815 0.08532309 0.07635770
##      8      19     123
## 0.06751845 0.06531750 0.05162090
```

```
## Plotting cook's distance
plot(mlr_sub1_model$cooks)
```



```
## Some observations have high cook's distance relative to other observations, but none have cook's d >
```


```
## Checking Constant Variance  
plot(mlr_sub1_model, which=1)
```



```
bptest(mlr_sub1_model)
```


```
##
## studentized Breusch-Pagan test
##
## data:  mlr_sub1_model
## BP = 77.914, df = 11, p-value = 3.73e-12
```

```
## Constant Variance seems to be violated
## Checking Normality
plot(mlr_sub1_model, which=2)
```



```
hist(mlr_sub1_model$residuals)
```

Histogram of mlr_sub1_model\$residuals



```
### We can use the KS test to assess normality because n>50.
ks.test(mlr_sub1_model$residuals, 'pnorm') ## We may want to check that this is the right syntax for t
```

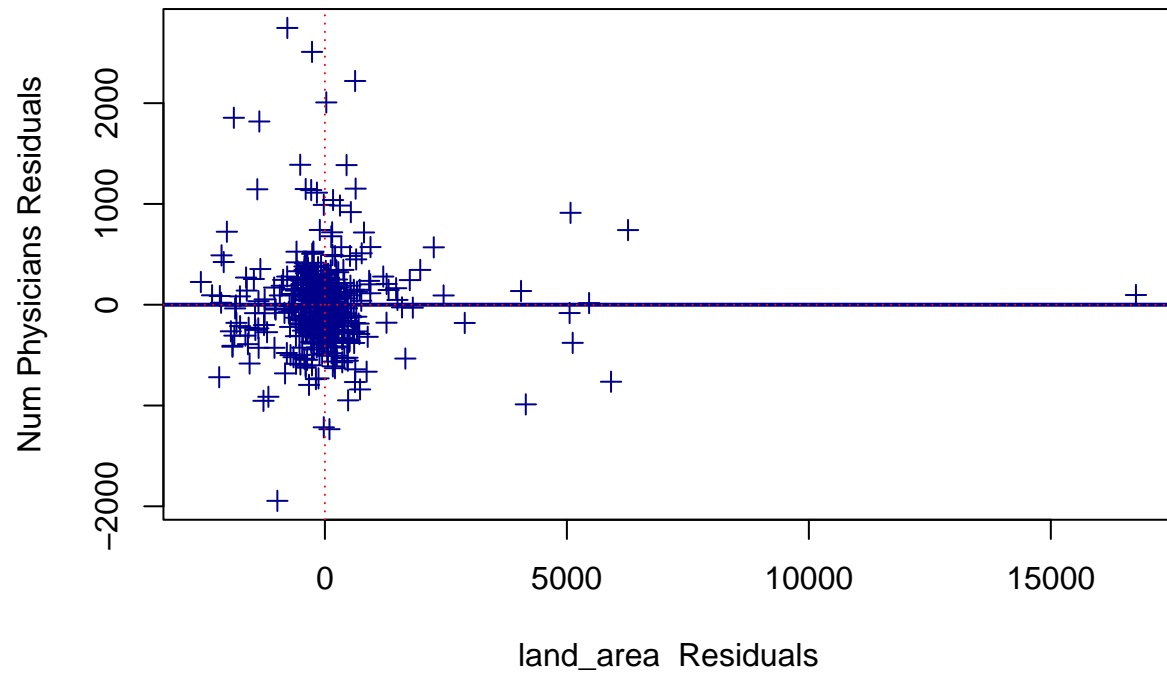
```
##
## Asymptotic one-sample Kolmogorov-Smirnov test
##
## data: mlr_sub1_model$residuals
## D = 0.52223, p-value < 2.2e-16
## alternative hypothesis: two-sided
```

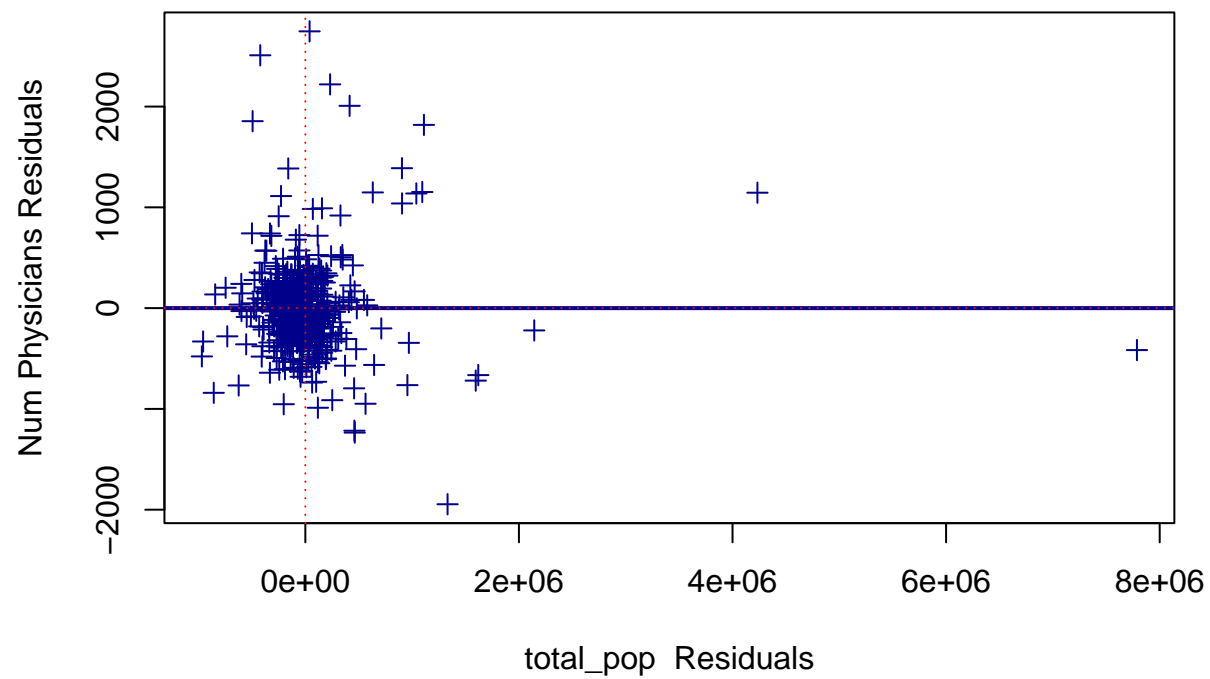
```
## Next step is to check linearity of each variable
```


```
checkLinearity <- function(var) {
  var_idx = which( colnames(sub1_data)==var )
  y.var = update(mlr_sub1_model, .~. -c(var_idx))$res
  #remove response + the variable itself
  x.var = lm(sub1_data[,var_idx] ~ . ,sub1_data[,-c(var_idx,4)])$res

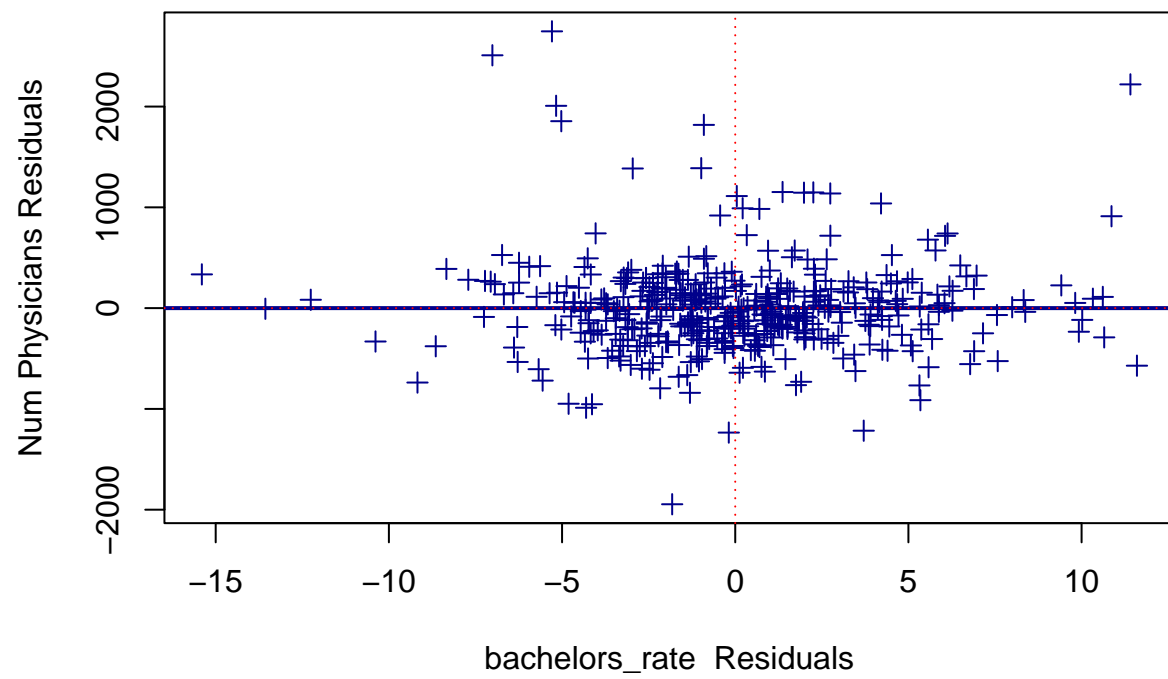
  plot(x.var, y.var, xlab=paste(var," Residuals"), ylab="Num Physicians Residuals", col='Darkblue', p
  abline(lm(y.var ~ x.var), col='Darkblue', lwd=2,xlim = c(quantile(x.var,.005),quantile(x.var,.995)))
  abline(v = 0, col="red", lty=3)
  abline(h = 0, col="red", lty=3)
}
predictors = names(sub1_data)
#remove the response variable (and region since it's a factor (?))
```


```
predictors = predictors[!(predictors %in% c("num_physicians","region"))]  
#check linearity for each predictor  
for (var in predictors) {  
  checkLinearity(var)  
}
```




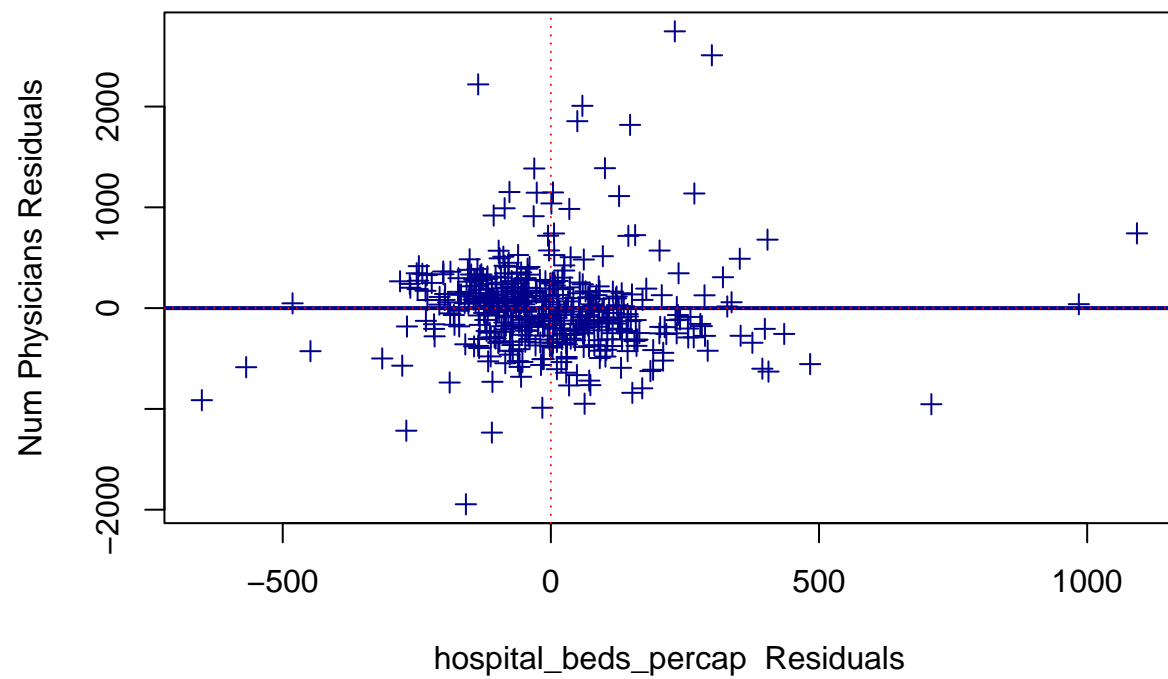
















```
# log transform target - Didn't help
sub2_data <- sub1_data %>%
  mutate(log_physicians = log(num_physicians))
mlr_sub2_model <- lm(log_physicians~., data = sub2_data)

# Checking Box Cox
physician.transformation = boxcox(mlr_sub1_model, lambda=seq(-2,2, length=400))
```




```
lambda <- physician.transformation$x[which.max(physician.transformation$y)]
lambda
```

```
## [1] 0.1152882
```

```
# Using 0.1 for box cox - Didn't help
lambda <- 0.1
sub2_data <- sub1_data %>%
  mutate(boxcox_physicians = (num_physicians^lambda - 1)/ lambda)
mlr_sub2_model <- lm(boxcox_physicians~., data = sub2_data)
```


```
## Checking Constant Variance
plot(mlr_sub2_model, which=1)
```

```
library(lmtest)
bptest(mlr_sub2_model)
```


```
##
## studentized Breusch-Pagan test
##
## data:  mlr_sub2_model
## BP = 231.29, df = 12, p-value < 2.2e-16
```

```
## Constant Variance seems to be violated
## Checking Normality
plot(mlr_sub2_model, which=2)
```



```
hist(mlr_sub2_model$residuals)
```


Histogram of mlr_sub2_model\$residuals




```
### We can use the KS test to assess normality because n>50.  
ks.test(mlr_sub2_model$residuals, 'pnorm')
```

```
##  
## Asymptotic one-sample Kolmogorov-Smirnov test  
##  
## data: mlr_sub2_model$residuals  
## D = 0.055808, p-value = 0.129  
## alternative hypothesis: two-sided
```


```
# Let's look at each variable and the target graphically  
# Need a log transformation on land_area  
model_data %>% ggplot(aes(x = log(land_area), y = log(num_physicians))) +  
  geom_point()
```




```
# log of total_pop will do wonders!!!  
model_data %>% ggplot(aes(x = log(total_pop), y = log(num_physicians))) +  
  geom_point()
```




```
# looks ok as is  
model_data %>% ggplot(aes(x = pop_18to24, y = log(num_physicians))) +  
  geom_point()
```




```
# looks ok as is  
model_data %>% ggplot(aes(x = pop_over65, y = log(num_physicians))) +  
  geom_point()
```




```
# looks ok as is - tried 1/x, x + x^2, sqrt and log but non look better  
model_data %>% ggplot(aes(x = highschool_rate, y = log(num_physicians))) +  
  geom_point()
```




```
# looks ok as is but log looks better  
model_data %>% ggplot(aes(x = log(bachelors_rate), y = log(num_physicians))) +  
  geom_point()
```


```
# looks ok as is  
model_data %>% ggplot(aes(x = poverty_rate, y = log(num_physicians))) +  
  geom_point()
```




```
# looks ok as is  
model_data %>% ggplot(aes(x = unemployment_rate, y = log(num_physicians))) +  
  geom_point()
```




```
# looks ok as is but log is better  
model_data %>% ggplot(aes(x = log(per_capita_income), y = log(num_physicians))) +  
  geom_point()
```



```
# log helps!  
model_data %>% ggplot(aes(x = log(hospital_beds_percap), y = log(num_physicians))) +  
  geom_point()
```



```
# looks ok as is but log is better  
model_data %>% ggplot(aes(x = log(serious_crimes_percap), y = log(num_physicians))) +  
  geom_point()
```



```
# log transformations based on above EDA
transformed_data <- model_data %>%
  mutate(log_physicians = log(num_physicians),
         log_pop = log(total_pop),
         log_land = log(land_area),
         log_bachelors = log(bachelors_rate),
         log_percap_income = log(per_capita_income),
         log_hospital_beds = log(hospital_beds_percap),
         log_crimes = log(serious_crimes_percap)) %>%
  dplyr::select(-c(num_physicians, total_pop, land_area, bachelors_rate, per_capita_income,
                  hospital_beds_percap, serious_crimes_percap))
```

```
# Start by just looking at total_pop for fun
slr_model <- lm(log_physicians~log_pop, data = transformed_data)
summary(slr_model)
```

```
##
## Call:
## lm(formula = log_physicians ~ log_pop, data = transformed_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.48926 -0.33040 -0.04263  0.27043  2.52197
##
## Coefficients:
```

```
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept) -10.06656    0.38025  -26.47  <2e-16 ***
## log_pop      1.29996    0.03042   42.74  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5037 on 438 degrees of freedom
## Multiple R-squared:  0.8066, Adjusted R-squared:  0.8061
## F-statistic: 1826 on 1 and 438 DF,  p-value: < 2.2e-16
```

Full Model

```
mlr_full_model <- lm(log_physicians~., data = transformed_data)
summary(mlr_full_model)
```

```
##
## Call:
## lm(formula = log_physicians ~ ., data = transformed_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.94655 -0.16551 -0.01646  0.14326  1.37012
##
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -19.899033    1.433753  -13.879  < 2e-16 ***
## pop_18to24      0.013936    0.004977   2.800  0.005348 **
## pop_over65      0.015032    0.004772   3.150  0.001748 **
## highschool_rate -0.002384    0.004195  -0.568  0.570094
## poverty_rate    0.023823    0.006340   3.757  0.000196 ***
## unemployment_rate -0.021619    0.008181  -2.643  0.008527 **
## region2        -0.096126    0.042693  -2.252  0.024858 *
## region3        -0.051341    0.045770  -1.122  0.262620
## region4         0.161340    0.054295   2.972  0.003131 **
## log_pop         1.072263    0.021622  49.591  < 2e-16 ***
## log_land       -0.044456    0.019109  -2.326  0.020466 *
## log_bachelors    0.569723    0.097012   5.873  8.66e-09 ***
## log_percap_income 0.771188    0.152647   5.052  6.50e-07 ***
## log_hospital_beds 0.543203    0.033362  16.282  < 2e-16 ***
## log_crimes      0.014359    0.037690   0.381  0.703406
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2689 on 425 degrees of freedom
## Multiple R-squared:  0.9465, Adjusted R-squared:  0.9447
## F-statistic: 537.1 on 14 and 425 DF,  p-value: < 2.2e-16
```

```
# Remove some fields based on EDA and full model summary
# The final group here was chosen based on EDA, p-values, and trial and error
sub1_data <- transformed_data %>%
  dplyr::select(-highschool_rate, -log_crimes, -log_land)
mlr_sub1_model <- lm(log_physicians~., data = sub1_data)
summary(mlr_sub1_model)
```

```
##
```

```
## Call:
## lm(formula = log_physicians ~ ., data = sub1_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.90134 -0.18275 -0.01244  0.14497  1.34825
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -21.634618    1.165895  -18.556 < 2e-16 ***
## pop_18to24      0.016813    0.004822   3.486 0.00054 ***
## pop_over65      0.015297    0.004784   3.198 0.00149 **
## poverty_rate    0.028214    0.005233   5.392 1.16e-07 ***
## unemployment_rate -0.023414    0.008070  -2.901 0.00391 **
## region2        -0.098454    0.039829  -2.472 0.01383 *
## region3        -0.042164    0.040105  -1.051 0.29370
## region4         0.120411    0.045368   2.654 0.00825 **
## log_pop         1.066383    0.019930  53.506 < 2e-16 ***
## log_bachelors    0.507090    0.075970   6.675 7.69e-11 ***
## log_percap_income 0.925893    0.137381   6.740 5.15e-11 ***
## log_hospital_beds 0.546572    0.032444  16.846 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2699 on 428 degrees of freedom
## Multiple R-squared:  0.9457, Adjusted R-squared:  0.9443
## F-statistic: 677.9 on 11 and 428 DF,  p-value: < 2.2e-16
```

```
anova(mlr_sub1_model, mlr_full_model)
```

```
## Analysis of Variance Table
##
## Model 1: log_physicians ~ pop_18to24 + pop_over65 + poverty_rate + unemployment_rate +
##      region + log_pop + log_bachelors + log_percap_income + log_hospital_beds
## Model 2: log_physicians ~ pop_18to24 + pop_over65 + highschool_rate +
##      poverty_rate + unemployment_rate + region + log_pop + log_land +
##      log_bachelors + log_percap_income + log_hospital_beds + log_crimes
##   Res.Df    RSS Df Sum of Sq    F Pr(>F)
## 1      428 31.189
## 2      425 30.740   3    0.44894 2.069 0.1037
```

```
# Adding a test for slr model vs full to show the the slr is rejected
anova(slr_model, mlr_full_model)
```

```
## Analysis of Variance Table
##
## Model 1: log_physicians ~ log_pop
## Model 2: log_physicians ~ pop_18to24 + pop_over65 + highschool_rate +
##      poverty_rate + unemployment_rate + region + log_pop + log_land +
##      log_bachelors + log_percap_income + log_hospital_beds + log_crimes
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      438 111.14
## 2      425  30.74  13    80.397 85.504 < 2.2e-16 ***
```




```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
## Checking high-leverage points
leverages=lm.influence(mlr_sub1_model)$hat
head(leverages)
```

```
##          1          2          3          4          5          6
## 0.06110336 0.04819432 0.03643543 0.03598304 0.03700481 0.05271986
```

```
## Plot to help identify high leverage observations
halfnorm(leverages, nlab=6, labs=as.character(1:length(leverages)), ylab="Leverages")
```



```
## Determining leverages that exceed a 2p/n threshold
n = dim(model_data)[1]
p = length(variable.names(mlr_sub1_model))
leverages.high = leverages[leverages>(2*p/n)]
leverages.high
```

```
##          1          42          67          95          128          155          171
## 0.06110336 0.12159176 0.05817187 0.07467461 0.14776376 0.06307226 0.08918100
##          173          184          187          188          206          246          249
## 0.08529432 0.05673314 0.06256259 0.09986278 0.07043394 0.05927733 0.06031901
##          272          301          303          334          337          357          363
```

```
## 0.05584776 0.06875972 0.13022693 0.05784573 0.09888911 0.06753811 0.09138324
##          396          398          404          405          412          415          433
## 0.07866738 0.09365475 0.12777068 0.07597522 0.08168668 0.09117746 0.06866181
##          436          437
## 0.07431743 0.09017867
```

```
## We currently have many high leverage points (30)
## Before continuing, let us look at what high leverage points are good and bad
## Calculate IQR for number of physicians
IQR_y = IQR(model_data$num_physicians)
## Define range with its lower limit being (Q1 - IQR) and upper limit being (Q3 + IQR)
QT1_y = quantile(model_data$num_physicians,0.25)
QT3_y = quantile(model_data$num_physicians,0.75)
lower_lim_y = QT1_y - IQR_y
upper_lim_y = QT3_y + IQR_y
vector_lim_y = c(lower_lim_y,upper_lim_y)
## Range for number of physicians
vector_lim_y
```

```
##      25%      75%
## -670.50 1889.25
```

```
## Extract observations with high leverage points from the original data frame
highlev = data[leverages>2*p/n,]
## Select only the observations with leverage points outside the range
highlev_lower = highlev[highlev$num_physicians < vector_lim_y[1], ]
highlev_upper = highlev[highlev$num_physicians > vector_lim_y[2], ]
highlev2 = rbind(highlev_lower,highlev_upper)

# Only 3 bad high leverage points
highlev2
```

```
## # A tibble: 3 x 17
##   id county      state land_~1 total~2 pop_1~3 pop_o~4 num_p~5 num_h~6 serio~7
##   <dbl> <chr>      <chr>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1    1 Los_Angel~ CA      4060 8863164   32.1    9.7   23677   27700  688936
## 2    67 Suffolk MA         59 663906   39.2   12.1    5674    6154   68808
## 3    95 Orleans LA        181 496938   28.3    13     2500    4018   54238
## # ... with 7 more variables: highschool_rate <dbl>, bachelors_rate <dbl>,
## # poverty_rate <dbl>, unemployment_rate <dbl>, per_capita_income <dbl>,
## # total_personal_income <dbl>, region <fct>, and abbreviated variable names
## # 1: land_area, 2: total_pop, 3: pop_18to24, 4: pop_over65,
## # 5: num_physicians, 6: num_hospital_beds, 7: serious_crimes
```

```
## Computing Studentized Residuals
mlr_sub1_model$resid = rstudent(mlr_sub1_model);

## Critical value with Bonferroni correction
## Note: Compare to t-value later at the alpha we choose
bonferroni_cv = qt(.05/(2*n), n-p-1)
bonferroni_cv
```

```
## [1] -3.895681
```

```
## Sorting residuals to find outliers
```

```
mlr_sub1_model.resid.sorted = sort(abs(mlr_sub1_model.resid), decreasing=TRUE)[1:10]  
print(mlr_sub1_model.resid.sorted)
```

```
##      418      42      431      291      258      50      248      282  
## 5.285196 4.004943 3.439459 3.369727 3.266127 2.922149 2.750065 2.687347  
##      346      241  
## 2.608512 2.548376
```

```
## 2 points are outliers (418, 42)
```

```
mlr_sub1_model.outliers = mlr_sub1_model.resid.sorted[abs(mlr_sub1_model.resid.sorted) > abs(bonferroni.  
print(mlr_sub1_model.outliers)
```

```
##      418      42  
## 5.285196 4.004943
```


```
## Finding high cook's distance observations
```

```
mlr_sub1_model.cooks = cooks.distance(mlr_sub1_model)  
sort(mlr_sub1_model.cooks, decreasing = TRUE)[1:10]
```

```
##      42      418      431      363      415      248      123  
## 0.17873948 0.11716642 0.03358325 0.02838177 0.02564037 0.02525072 0.02435829  
##      181      258      406  
## 0.02274234 0.02142160 0.01913981
```


```
## Plotting cook's distance
```

```
plot(mlr_sub1_model.cooks)
```



```
## Some observations have high cook's distance relative to other observations, but none have cook's d >
```


```
## Checking Constant Variance  
plot(mlr_sub1_model, which=1)
```



```
bptest(mlr_sub1_model)
```


```
##
## studentized Breusch-Pagan test
##
## data:  mlr_sub1_model
## BP = 37.381, df = 11, p-value = 9.945e-05
```

```
## Constant Variance seems to be violated based on p-value, the plot looks OK though
## Checking Normality
plot(mlr_sub1_model, which=2)
```



```
hist(mlr_sub1_model$residuals)
```

Histogram of mlr_sub1_model\$residuals



```
### We can use the KS test to assess normality because n>50.
ks.test(mlr_sub1_model$residuals, 'pnorm') ## We may want to check that this is the right syntax for t
```

```
##
## Asymptotic one-sample Kolmogorov-Smirnov test
##
## data: mlr_sub1_model$residuals
## D = 0.30228, p-value < 2.2e-16
## alternative hypothesis: two-sided
```

```
## Next step is to check linearity of each variable
```

```
sub1_data <- data.frame(sub1_data)


checkLinearity <- function(var) {
  var_idx = which( colnames(sub1_data)==var )
  y.var = update(mlr_sub1_model, .~. -c(var_idx))$res
  #remove response + the variable itself
  x.var = lm(sub1_data[,var_idx] ~ . ,sub1_data[,-c(6)])$res


  plot(x.var, y.var, xlab=paste(var," Residuals"), ylab="Num Physicians Residuals", col='Darkblue', p
  abline(lm(y.var ~ x.var), col='Darkblue', lwd=2)
  abline(v = 0, col="red", lty=3)
  abline(h = 0, col="red", lty=3)
}
```


```


predictors = names(sub1_data)
#remove the response variable (and region since it's a factor (?))
predictors = predictors[!(predictors %in% c("log_physicians","region"))]
#check linearity for each predictor
for (var in predictors) {
  checkLinearity(var)
}


```




















```
## Checking VIF
```

```
round(vif(mlr_sub1_model),3) #log_bachelors and log_percap_income are kind of concerning to me, but are
```

```
##      pop_18to24      pop_over65      poverty_rate      unemployment_rate
##      2.461      2.198      3.577      2.144
##      region2      region3      region4      log_pop
##      1.774      2.196      1.794      1.495
##      log_bachelors      log_percap_income      log_hospital_beds
##      4.368      4.859      1.924
```

```
## Grabbing design matrix
```

```
x = model.matrix(mlr_sub1_model)[,-1]
```

```
## Standardize the matrix
```

```
x = x - matrix(apply(x,2, mean), 440,11, byrow=TRUE)
```

```
x = x / matrix(apply(x, 2, sd), 440,11, byrow=TRUE)
```

```
## Compute the eigenvalues of the matrix
```

```
eigenvalues.x = eigen(t(x) %*% x)
```

```
eigenvalues.x$val
```

```
## [1] 1275.27414 803.29800 668.63050 624.17316 530.54827 321.79620
## [7] 219.22558 147.41423 111.62601 84.01384 43.00007
```

```
## Compute Condition Number  
sqrt(eigenvalues.x$val[1]/eigenvalues.x$val[8]) ## Is less than 30, looks good.
```

```
## [1] 2.941251
```